



Figure 1: Certification

DECLARATION

I **NIREEKSHITH B R** [4HG21CS028] student of 6th semester B.E, CS & E, Government Engineering College Mosalehosahalli, Hassan, hereby declare that the internship entitled “**Internship Training Programme**” has been carried out by me through the **Internshala Trainings**. This internship report is submitted in partial fulfillment of the requirements for the award of the degree of Computer Science and Engineering by Visvesvaraya Technological University during the academic year 2023-24. This report has not been submitted to any other organization/university for any award of a degree certificate.

BY:

NIREEKSHITH B R[4HG21CS028]

ABSTRACT

The implementation of cloud computing using Amazon Web Services (AWS) within a college network infrastructure presents a comprehensive approach to scalable, flexible, and cost-effective resource management. AWS provides a wide array of cloud services, including computing power, storage options, and networking capabilities, that enable educational institutions to enhance their IT infrastructure significantly.

This internship training explores the implementation of cloud computing using Amazon Web Services (AWS), aiming to enhance scalability, flexibility, and cost-efficiency within an educational context. By utilizing AWS core services such as Amazon EC2 for computing power, Amazon S3 for storage, Amazon RDS for database management, and AWS Lambda for serverless computing, the training demonstrates practical applications for deploying and managing web applications, databases, and virtual machines.

Key security measures are emphasized through AWS Identity and Access Management (IAM) for secure access control and AWS Virtual Private Cloud (VPC) for creating isolated network environments, ensuring data protection and secure communication. The training leverages AWS's robust tools to provide hands-on learning experiences, offering valuable insights into cloud architecture, deployment strategies, and best practices for optimizing performance and cost.

Through this internship, the seamless integration of AWS cloud computing into educational IT infrastructure is illustrated, facilitating efficient data management, scalability, and accessibility across different departments and user groups. This training underscores the transformative potential of cloud technology in modernizing educational institutions, enhancing their IT resource management and operational efficiency.

ACKNOWLEDGEMENT

The completion of any Intrenship involves the efforts of many people. I have been lucky to have received a lot of help, support from all directions during this Intrenship Program, so with the gratitude i take this opportunity to acknowledge all those who guide and encouraged me.

I would like to express our profound sense of gratitude to coordinator **Dr.Raghu M.E**, Assistant Professor, Dept. of CS & E , **Dr.K C Ravishankar**, Professor and HOD, Dept. of CS & E and our Principal **Dr.T Rangaswamy**, creating an excellent and technically sound enviornment in our institustion.It's a great pleasure to acknowledge for their help and encouragement rendered towards the successful completion of Intrenship.

BY:

NIREEKSHITH B R [4HG21CS028]

Contents

1	Introduction	1
1.1	Introduction on Internshala Trainings	1
1.1.1	Vision:	2
1.1.2	Mission:	2
1.2	Cloud Computing	2
1.3	Amazon Web Services	3
1.4	Objectives	3
1.5	Application	4
1.6	Advantages	4
1.7	Disadvantages	4
1.8	Organization of Report	5
2	AWS Service Overview	7
3	Requirement Analysis	12
3.1	Hardware Requirements	12
3.2	Software Requirements	12
3.3	Functional Requirements	13
3.4	Non Functional Requirements	13
4	AWS Deployment Models	15
4.1	Deployment Models	15
5	AWS Architecture	17
6	AWS Pricing and Cost Management	20

7 Case Studies and Use Cases	22
8 Future Trends in Cloud Computing and AWS	25
9 Conclusion	27
Bibliography	28

List of Figures

1	Certification	
1.1	Internshala Interface	1
2.1	AWS EC2 Instance	7
2.2	Simple Storage Bucket	8
2.3	Virtual Private Cloud	9
2.4	Relational Database Service	9
2.5	AWS Security Group	10
2.6	AWS Management Console	11
4.1	AWS Services	15
5.1	AWS Architecture	17
6.1	AWS Cost Management	20
7.1	AWS Case study and Use Cases	22

List of Tables

3.1	Minimum Hardware Requirement	12
3.2	Minimum Software Requirement	12

Chapter 1

Introduction

1.1 Introduction on Internshala Trainings

Internshala Trainings provides a wide range of online training programs tailored for students and professionals seeking to enhance their skills and career prospects. Whether you're interested in programming, digital marketing, graphic design, or other fields, Internshala Trainings offers courses to suit various interests and career goals. Upon completion of these courses, participants receive certificates, adding value to their resumes and LinkedIn profiles.

The curriculum of Internshala Trainings is carefully crafted by industry experts to provide hands-on experience and industry-relevant skills. Participants not only learn theoretical concepts but also gain practical knowledge that can be applied in real-world scenarios. The courses cover diverse areas such as programming languages, web development, data science, marketing, finance, creative arts, and more.




Figure 1.1: Internshala Interface

Internshala Trainings offers flexible learning options, allowing participants to study at their own pace and convenience from anywhere with an internet connection. Moreover, these courses are priced affordably, making quality education accessible to everyone. Additionally, completion of certain courses may enhance internship opportunities through Internshala's platform, providing a comprehensive learning and career growth experience.

1.1.1 Vision:

The vision of Internshala Trainings is to democratize education by providing high-quality, affordable, and accessible training programs. We aim to empower individuals worldwide, regardless of their background or location, by offering courses that bridge the gap between theoretical knowledge and practical skills, thus enabling them to succeed in their careers and achieve their professional goals.

1.1.2 Mission:

Internshala Trainings' mission is to equip individuals with the skills they need to succeed in the ever-evolving professional world. We strive to achieve this by offering industry-relevant training programs that are affordable, accessible, and practical. Our mission is to empower students and professionals worldwide to learn new skills, advance in their careers, and fulfill their aspirations through quality education and practical training experiences.

1.2 Cloud Computing

Cloud computing revolutionizes the way IT resources are provisioned, managed, and utilized. At its core, cloud computing refers to the delivery of computing services—including servers, storage, databases, networking, software, and more—over the internet, providing users with on-demand access to a shared pool of configurable resources. Unlike traditional computing models that require users to own and maintain physical hardware and infrastructure, cloud computing allows organizations to access resources as a utility, paying only for what they use.

Cloud computing offers several key benefits, including scalability, flexibility, cost-efficiency, and reliability. With cloud services, organizations can rapidly scale their resources up or down

based on demand, allowing for agility in responding to changing business needs. Additionally, cloud computing enables users to access their applications and data from anywhere with an internet connection, promoting flexibility and remote work. Moreover, by shifting from capital expenditure (CapEx) to operational expenditure (OpEx), organizations can reduce costs and only pay for the resources they consume. Cloud providers ensure high reliability and availability of services through redundant infrastructure and data centers.

1.3 Amazon Web Services

Amazon Web Services (AWS) is a leading cloud service provider that offers a vast array of on-demand computing resources and services over the internet. Launched in 2006, AWS provides a reliable, scalable, and secure platform for businesses, governments, and individuals to build and deploy applications, store data, and utilize various IT resources without the need for owning physical infrastructure.

AWS offers a wide range of services across compute, storage, databases, networking, machine learning, security, analytics, and more. Some core services include Amazon Elastic Compute Cloud (EC2) for scalable virtual servers, Amazon Simple Storage Service (S3) for object storage, Amazon Relational Database Service (RDS) for managed databases, Amazon Virtual Private Cloud (VPC) for networking, and Amazon Simple Queue Service (SQS) for message queuing.

1.4 Objectives

1. To provide scalable computing resources that can easily scale up or down based on demand, allowing businesses to accommodate fluctuating workloads efficiently.
2. To offer highly reliable and available services with redundant infrastructure and data centers, ensuring minimal downtime and uninterrupted operation.
3. To enable cost-effective IT solutions by allowing users to pay only for the resources they use, without upfront investment in hardware or long-term contracts.
4. To provide a wide range of services and tools that cater to various business needs, allowing users to build, deploy, and manage applications with flexibility.

1.5 Application

- **Web Hosting and Applications:** Hosting websites, web applications, and APIs on scalable infrastructure using services like EC2, S3, and AWS Lambda.
- **Data Storage and Backup:** Storing and backing up data securely using services such as Amazon S3, Glacier, and EBS.
- **Big Data Analytics:** Processing and analyzing large datasets using services like Amazon Redshift, EMR (Elastic MapReduce), and Athena.
- **Machine Learning and AI:** Building and deploying machine learning models with services like Amazon SageMaker, Rekognition, and Comprehend.
- **DevOps and Continuous Integration/Continuous Deployment (CI/CD):** Automating software delivery pipelines using services such as AWS CodePipeline, CodeBuild, and CodeDeploy.

1.6 Advantages

- **Scalability:** Easily scale resources up or down to handle changing workloads.
- **Reliability:** High availability and reliability with redundant infrastructure and data centers.
- **Cost Efficiency:** Pay-as-you-go pricing model helps in cost optimization by paying only for what you use.
- **Global Reach:** Global infrastructure with data centers in multiple regions enables low-latency access worldwide.
- **Wide Range of Services:** Comprehensive suite of services for various use cases, from computing and storage to AI and IoT.

1.7 Disadvantages

- **Complexity:** Managing AWS resources can be complex, especially for beginners, due to the wide range of services and configurations available.

- **Cost Management:** While pay-as-you-go can be cost-effective, it requires careful monitoring and management to avoid unexpected bills.
- **Vendor Lock-in:** Using AWS-specific services may result in vendor lock-in, making it challenging to migrate to other platforms in the future.
- **Security Concerns:** Security breaches or data leaks can occur if not properly configured or secured.
- **Dependency on Internet Connectivity:** Relies on internet connectivity; downtime or network issues can impact access to services.

1.8 Organization of Report

This internship in web development with a project is divided up into chapters, each dealing with different aspects of the amazon web services. Each chapter has a short introduction explaining the subject of each chapter, and then the details Each module is explained separately. The following is a short overview of each of the chapters

1. Introduction to Cloud Computing:

In this chapter, we discussed the fundamentals of cloud computing, including its definition, importance, and key benefits. Cloud service models (IaaS, PaaS, SaaS) and deployment models (public, private, hybrid) were also introduced.

2. Setting Started with AWS:

This chapter covers the basics of getting started with Amazon Web Services (AWS), including signing up for an AWS account, navigating the AWS Management Console, and setting up security credentials.

3. AWS Identity and Access Management (IAM):

IAM allows you to manage access to AWS services securely. This chapter discusses IAM concepts, such as users, groups, roles, and policies, and how to control access to AWS resources.

4. Simple Storage Service (S3):

S3 is a scalable object storage service provided by AWS. Here, we explore how to create S3 buckets, upload and manage objects, set permissions, and enable versioning.

5. Virtual Private Cloud (VPC):

VPC enables you to create isolated virtual networks within the AWS cloud. This chapter covers VPC concepts, subnets, route tables, security groups, and network ACLs.

6. Elastic Compute Cloud (EC2):

EC2 provides scalable virtual servers in the cloud. Here, we discuss launching EC2 instances, choosing instance types, configuring security, and connecting to instances.

7. Higher-Level Architecture:

This chapter discusses designing higher-level architectures on AWS, including load balancing, auto-scaling, content delivery, and fault tolerance using services like ELB, Auto Scaling, CloudFront, and Route 53.

8. Database Services:

AWS offers various database services like RDS, DynamoDB, and Aurora. This chapter covers setting up and managing databases, data replication, backups, and scaling.

9. AWS Security and Management Services:

Here, we explore additional AWS security services such as AWS WAF, Shield, Inspector, and CloudTrail, along with management services like CloudFormation and AWS Config.

10. AWS Billing and Pricing:

This chapter explains AWS billing models, cost management strategies, and tools for monitoring and optimizing costs, such as AWS Cost Explorer and Budgets.

11. Training Conclusion and Final Project:

In this concluding chapter, we summarize key learnings from the training. Participants will also work on a final project to apply the knowledge gained throughout the training.

Chapter 2

AWS Service Overview

Compute:

1. EC2 (Elastic Compute Cloud):

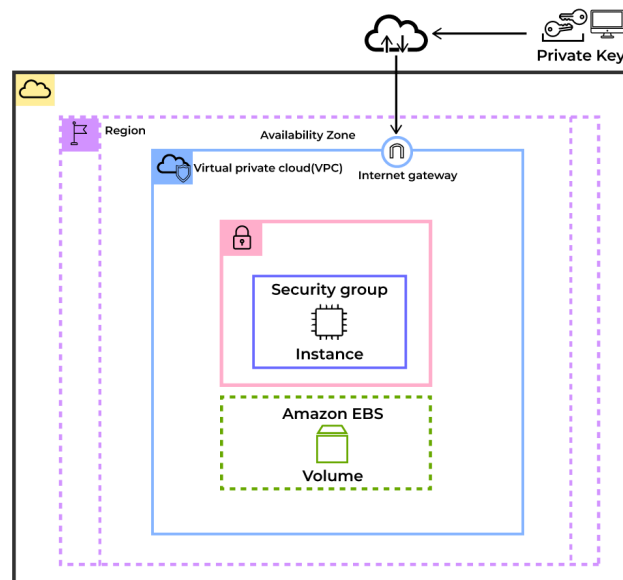


Figure 2.1: AWS EC2 Instance

Amazon EC2 provides resizable compute capacity in the cloud, allowing you to launch virtual servers known as instances. You can choose instance types based on your application's requirements, such as CPU, memory, storage, and networking capacity. EC2 instances are scalable, flexible, and can be easily configured to run various operating systems and applications, making them suitable for a wide range of workloads from simple web hosting to complex data analysis tasks.

2.Lambda:

AWS Lambda is a serverless compute service that lets you run code in response to events without provisioning or managing servers. You can upload your code and Lambda automatically scales to handle the incoming traffic or events. It's ideal for executing small, event-driven functions that respond to changes in data, such as file uploads, database modifications, or incoming API requests. With Lambda, you pay only for the compute time consumed by your code.

Storage:

1. S3 (Simple Storage Service):

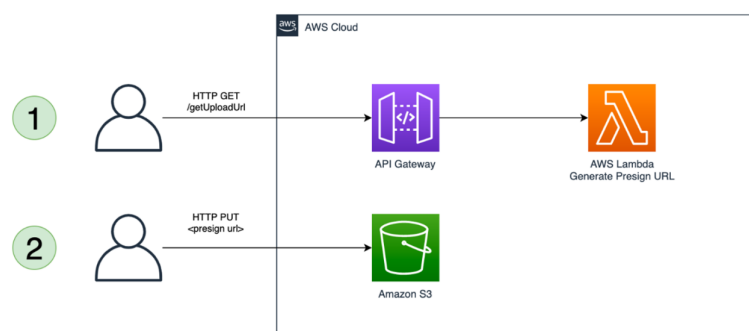


Figure 2.2: Simple Storage Bucket

Amazon S3 is an object storage service designed to store and retrieve any amount of data from anywhere on the web. It offers scalability, high availability, and durability. S3 allows you to store and retrieve data objects, such as photos, videos, documents, and backups. You can control access to your data, configure storage classes for cost optimization, and set up event notifications for automation.

2. EBS (Elastic Block Store):

Elastic Block Store provides block-level storage volumes that you can attach to EC2 instances to persist data beyond the life of the instance. EBS volumes are highly available and reliable, suitable for databases, application storage, and boot volumes. You can choose different volume types based on performance characteristics and use cases, such as SSD-backed volumes for high-performance workloads and HDD-backed volumes for throughput-intensive applications.

Networking:

1. VPC (Virtual Private Cloud):

Amazon VPC allows you to create a virtual network in the AWS cloud, providing isolation for your resources. Within a VPC, you can define subnets, route tables, and network gateways.

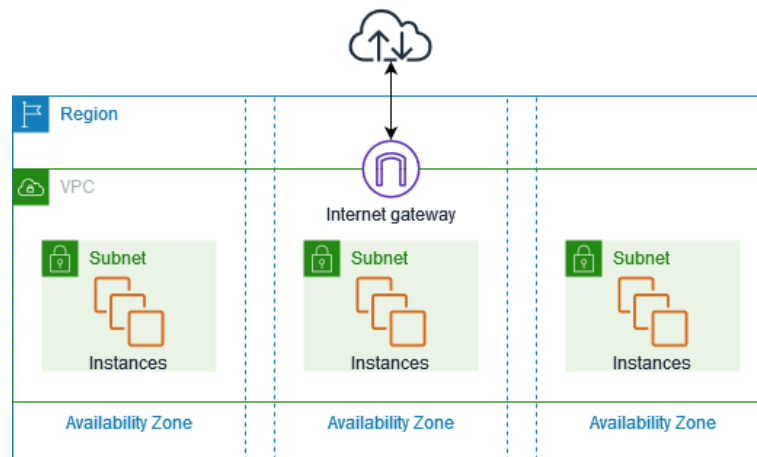


Figure 2.3: Virtual Private Cloud

It enables you to control your network settings, such as IP addressing, internet access, and communication between resources. VPC lets you extend your data center into the cloud while maintaining security and isolation.

2. Route 53:

Amazon Route 53 is a highly available and scalable DNS web service. It translates human-readable domain names into IP addresses and routes users to resources like EC2 instances, load balancers, or S3 buckets. Route 53 supports domain registration, health checks for routing traffic to healthy endpoints, and DNS failover to improve the availability and reliability of your applications.

Databases:

1. RDS (Relational Database Service):




Figure 2.4: Relational Database Service

Amazon RDS is a managed relational database service that simplifies database administration

tasks. It supports several database engines like MySQL, PostgreSQL, SQL Server, etc. RDS manages routine database tasks such as provisioning, patching, backup, recovery, and scaling, allowing you to focus on your applications rather than database management.

2.DynamoDB:

Amazon DynamoDB is a fully managed NoSQL database service designed for applications that require single-digit millisecond latency at any scale. It provides seamless scalability with consistent, fast performance. DynamoDB is suited for a wide range of applications, from gaming to IoT, and offers features like automatic scaling, backup and restore, and flexible data modeling.

Security Identity:

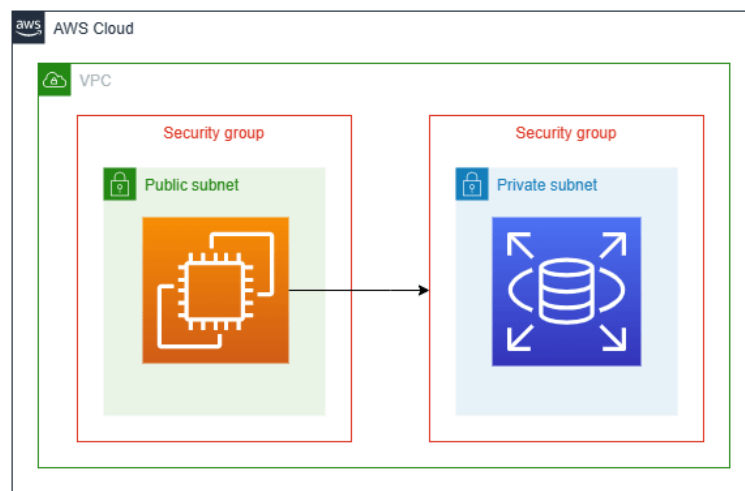


Figure 2.5: AWS Security Group

1. IAM (Identity and Access Management):

AWS Identity and Access Management (IAM) enables you to manage access to AWS services and resources securely. IAM allows you to create and manage users, groups, and roles and define granular permissions to control access to resources. It helps you implement the principle of least privilege and ensures that only authorized entities can interact with your AWS environment.

2.Security Groups:

Security Groups act as virtual firewalls for your EC2 instances, controlling inbound and outbound traffic. You can define rules that allow specific types of traffic to reach your instances, providing a layer of security at the instance level. Security Groups are stateful, meaning if you allow outbound traffic, the response traffic is automatically allowed, simplifying network security configuration.

Management Tools: Security Identity:

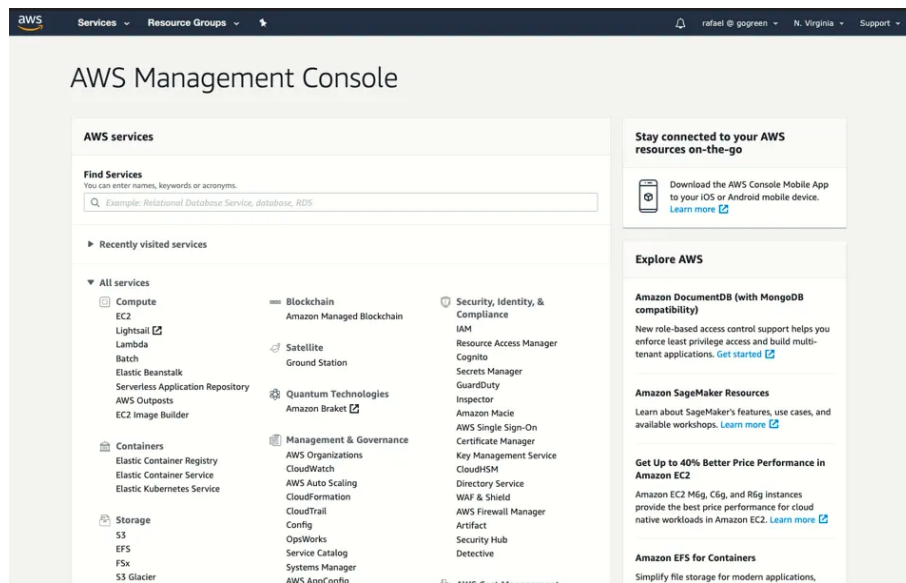


Figure 2.6: AWS Management Console

1.CloudWatch:

Amazon CloudWatch is a monitoring and observability service that provides data and actionable insights for AWS resources and applications. It collects and tracks metrics, monitors logs and events, sets alarms, and automatically reacts to changes in your AWS resources. CloudWatch helps you gain visibility into your applications' performance, resource utilization, and operational health.

2.AWS CLI (Command Line Interface):

The AWS Command Line Interface is a unified tool to manage AWS services from the command line. It allows you to control multiple AWS services directly from your terminal, automate tasks, and integrate with scripts and programs. The CLI provides easy access to AWS resources and is particularly useful for automation, scripting, and managing AWS resources in text-based environments.

Chapter 3

Requirement Analysis

3.1 Hardware Requirements

Name	Minimum Requirement
Processor	Intel/AMD (2Ghz)
RAM	4Gb
Hard Disk	10Gb

Table 3.1: Minimum Hardware Requirement

3.2 Software Requirements

Component	Requirement
Operating System	Windows 10 / Linux / MAC
Application Dependencies	specific libraries,frameworks
Database Software	mysql,postgresql,dynamDB etc
Development Tools	IDES,SDKS,compilers,version control systems
Monitoring Devices	AWS CLI,cloud watch,monitoring agents

Table 3.2: Minimum Software Requirement

Operating System: The operating system required to run the application software. It could be Linux, Windows, or another OS based on application compatibility and requirements.

Application Dependencies: Libraries, frameworks, or software required to run the applica-

tion code. This could include runtime environments like Java Runtime Environment (JRE), .NET Framework, etc.

Database Software: Software needed to manage databases. This might include MySQL, PostgreSQL, or other database engines depending on application needs.

Development Tools: Tools required for development, testing, and deployment, such as Integrated Development Environments (IDEs), compilers, and version control systems.

Monitoring and Management Tools: Description: Tools needed for monitoring and managing the system, such as AWS CLI, CloudWatch, monitoring agents, etc.

3.3 Functional Requirements

1. Compute Functionality:

Requirements related to compute resources and their functionality, such as deployment, management, and scaling.

2. Data Storage:

Requirements for storing and retrieving data reliably, including CRUD operations and data integrity.

3. Networking:

Requirements related to network configuration, DNS setup, and communication between components.

4. Database Functionality:

Requirements related to database operations such as CRUD operations, indexing, and querying.

3.4 Non Functional Requirements

1. Performance:

Requirements related to system performance, including response time, throughput, and scalability.

2. **Reliability:** Requirements related to system reliability, uptime, fault tolerance, and disaster recovery.

3. **Security:**

Requirements related to system security, including access control, encryption, and network security.

4. **Usability:**

Requirements related to system usability, including user interfaces and ease of use.

5. **Scalability:**

Requirements related to system scalability, including horizontal and vertical scalability options.

Chapter 4

AWS Deployment Models

AWS offers various deployment models catering to different needs. Organizations can choose on-premises deployment for maintaining applications in their own data centers with full control over hardware. In Infrastructure as a Service (IaaS), AWS provides virtualized computing resources over the internet, allowing users to manage applications while AWS handles infrastructure. Platform as a Service (PaaS) abstracts infrastructure management, enabling developers to focus solely on application development. AWS supports containerization with ECS and EKS for efficient application packaging and deployment. With serverless computing on Lambda, developers can deploy code without managing servers, paying only for actual usage.




Figure 4.1: AWS Services

4.1 Deployment Models

1. On-Premises Deployment:

In an on-premises deployment model, organizations deploy applications and infrastructure within their own data centers or private infrastructure. AWS provides services like AWS Outposts and VMware Cloud on AWS to extend AWS services to on-premises environments. Organizations retain full control over hardware, networking, and security but may lack the scalability and agility of the cloud.

2. **Infrastructure as a Service (IaaS):** In the IaaS model, AWS provides virtualized computing resources over the internet. Customers can rent virtual machines (EC2), storage (S3, EBS), and networking components (VPC) to build their own infrastructure. Customers manage applications, data, runtime, middleware, and operating systems while AWS manages the infrastructure.
3. **Platform as a Service (PaaS):** PaaS provides a platform allowing customers to develop, run, and manage applications without dealing with the underlying infrastructure. AWS services like AWS Elastic Beanstalk, AWS Lambda, and AWS Fargate fall under this category. Customers focus on application development and deployment while AWS handles infrastructure management, scaling, and maintenance.

4. **Containerization:**

Containerization allows packaging applications and their dependencies into containers for easy deployment and scalability. AWS provides services like Amazon ECS (Elastic Container Service) and Amazon EKS (Elastic Kubernetes Service) for container orchestration. Containers offer portability across environments and enable efficient resource utilization.

5. **Serverless Computing:** Serverless computing abstracts server management, allowing developers to focus solely on writing code. AWS Lambda is a key serverless compute service where users upload code and AWS handles the rest, including scaling and management. Serverless architectures can significantly reduce operational overhead and costs by charging only for actual usage.
6. **Hybrid Deployment:** Hybrid deployment combines on-premises infrastructure with cloud resources. AWS offers services like AWS Direct Connect and VPN to establish secure connections between on-premises data centers and AWS. Organizations can leverage the scalability and flexibility of the cloud while maintaining some applications or data on-premises for compliance or other reasons.

AWS Architecture




Figure 5.1: AWS Architecture

Follow AWS Well-Architected Framework pillars: Operational Excellence, Security, Reliability, Performance Efficiency, and Cost Optimization.

Design for scalability by leveraging AWS Auto Scaling, which automatically adjusts resources based on demand. Use services like Amazon EC2, ECS, or Lambda for scalable compute, and

Amazon RDS or DynamoDB for scalable databases.

High Availability and Fault Tolerance:

Design with redundancy across multiple Availability Zones (AZs) to ensure high availability.

Use load balancing and failover mechanisms (ELB, Route 53) for fault tolerance.

Security:

Implement least privilege access using IAM roles and policies. Encrypt data at rest and in transit using services like AWS KMS, SSL/TLS. Enable logging, monitoring, and auditing using AWS CloudTrail, CloudWatch, and AWS Config.

Resilient Networking:

Design VPC with public and private subnets for security and segmentation. Use Network ACLs and Security Groups to control inbound and outbound traffic. Implement VPN or AWS Direct Connect for secure connectivity.

Cost Optimization:

Right-size resources to match workload requirements and use AWS Cost Explorer to analyze costs. Utilize AWS Reserved Instances, Spot Instances, and Savings Plans for cost savings. Implement tagging and automation to manage costs effectively.

Automation and Infrastructure as Code (IaC):

Use AWS CloudFormation or AWS CDK for infrastructure provisioning and management. Implement configuration management with AWS Systems Manager (SSM) or third-party tools like Ansible, Chef, or Puppet.

Backup and Disaster Recovery:

Implement automated backups using services like Amazon S3, EBS Snapshots, or RDS automated backups. Design disaster recovery strategies with multi-region deployments or using AWS Backup.

Performance Optimization:

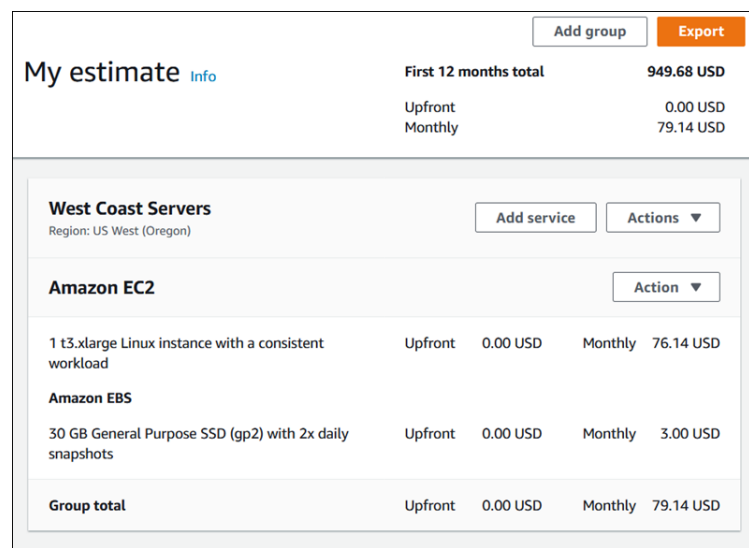
Optimize performance using AWS services like Amazon CloudFront for content delivery, Amazon ElastiCache for caching, and Amazon Aurora for high-performance databases. Monitor and optimize performance using AWS CloudWatch metrics and AWS Trusted Advisor.

Testing and Continuous Integration/Continuous Deployment (CI/CD):

Implement CI/CD pipelines using AWS CodePipeline, CodeBuild, and CodeDeploy. Automate testing and deployments to improve agility and reliability.

Chapter 6

AWS Pricing and Cost Management



The screenshot displays the AWS Pricing Calculator interface. At the top, there's a header with 'My estimate' and an 'Info' link. To the right, there are buttons for 'Add group' and 'Export'. Below this, a summary table shows the 'First 12 months total' cost as 949.68 USD, with 'Upfront' costs at 0.00 USD and 'Monthly' costs at 79.14 USD. The main section is titled 'West Coast Servers' with a sub-note 'Region: US West (Oregon)'. It includes buttons for 'Add service' and 'Actions'. Underneath, there's a section for 'Amazon EC2' with an 'Action' button. The table lists two services: '1 t3.xlarge Linux instance with a consistent workload' and '30 GB General Purpose SSD (gp2) with 2x daily snapshots'. Each service has columns for 'Upfront' and 'Monthly' costs. A 'Group total' row at the bottom summarizes the costs: Upfront 0.00 USD and Monthly 79.14 USD.

My estimate Info		First 12 months total	
Upfront	0.00 USD	Monthly	79.14 USD
Monthly	79.14 USD		

West Coast Servers		Add service		Actions ▼	
Region: US West (Oregon)					
Amazon EC2				Action ▼	
1 t3.xlarge Linux instance with a consistent workload	Upfront	0.00 USD	Monthly	76.14 USD	
Amazon EBS					
30 GB General Purpose SSD (gp2) with 2x daily snapshots	Upfront	0.00 USD	Monthly	3.00 USD	
Group total	Upfront	0.00 USD	Monthly	79.14 USD	

Figure 6.1: AWS Cost Management

Pay-As-You-Go Model:

AWS follows a pay-as-you-go pricing model where you pay only for the resources you consume. No upfront costs or long-term commitments are required for most services.

Pricing Calculator:

Use the AWS Pricing Calculator to estimate the costs of AWS services based on your usage patterns and requirements. It helps in planning and budgeting for AWS resources before deployment.

Cost Explorer:

AWS Cost Explorer provides insights into your AWS spending, allowing you to visualize, un-

derstand, and manage costs over time. Analyze costs by service, usage type, region, and more.

Cost Allocation Tags:

Use cost allocation tags to categorize and track spending across different projects, departments, or environments. Helps in understanding and optimizing costs at a granular level.

Reserved Instances (RIs):

RIs provide significant discounts (up to 75percent) compared to on-demand pricing in exchange for a commitment to usage over a 1 or 3-year term. Ideal for predictable workloads with steady-state usage.

Savings Plans:

Savings Plans offer savings on AWS usage in exchange for a commitment to a consistent amount of usage (measured in hour) over a 1 or 3-year term. Provide flexibility across EC2, Lambda, and Fargate usage.

Spot Instances and Spot Fleets:

Spot Instances allow you to bid on unused EC2 capacity, offering significant savings for fault-tolerant and flexible workloads. Spot Fleets allow you to manage a group of Spot Instances along with On-Demand and Reserved Instances for cost optimization.

Auto Scaling and Right Sizing:

Use AWS Auto Scaling to automatically adjust resources based on demand, optimizing costs. Right-size instances and resources based on actual usage to avoid over-provisioning.

AWS Budgets and Alerts:

Set up AWS Budgets to track spending and receive alerts when costs exceed predefined thresholds. Helps in proactively managing costs and avoiding unexpected bills.

Free Tier and Cost Optimization Recommendations:

Take advantage of the AWS Free Tier to explore and experiment with AWS services at no cost. Use AWS Cost Explorer's Cost Optimization Recommendations to identify potential cost-saving opportunities.

Managed Services and Serverless Computing:

Utilize managed services like AWS Lambda, S3, and DynamoDB where you pay only for what you use, reducing operational overhead and costs.

Chapter 7

Case Studies and Use Cases

The below examples illustrate the diverse applications and benefits of using AWS across various industries, from media streaming and travel to industrial IoT and space exploration. Each case demonstrates how AWS services can address specific challenges, enhance performance, and drive innovation.




Figure 7.1: AWS Case study and Use Cases

1. Netflix: Media Streaming

Challenge: Netflix needed a scalable infrastructure to support its global streaming service.

Solution: Migrated to AWS, using services like EC2 for compute, S3 for storage, and CloudFront for content delivery.

Outcome: Achieved high availability, scalability, and cost savings, allowing seamless streaming to millions of users worldwide.

2. Airbnb: Scalability and Reliability

Challenge: Airbnb required a robust and scalable infrastructure to handle peak loads and ensure reliability.

Solution: Implemented AWS services like EC2 for scalable compute resources, RDS for managed database services, and S3 for durable storage.

Outcome: Improved scalability, reliability, and performance, enabling rapid growth and global expansion.

3. GE: Industrial IoT

Challenge: GE needed a scalable platform for its Industrial Internet of Things (IIoT) applications.

Solution: Used AWS IoT, Lambda for serverless computing, and Redshift for data warehousing.

Outcome: Developed a scalable IIoT platform, enhancing data analytics capabilities and operational efficiency.

4. Expedia: Big Data Analytics

Challenge: Expedia needed to analyze large volumes of travel data to improve customer experiences.

Solution: Leveraged AWS services like EMR for big data processing, S3 for data storage, and Redshift for data warehousing.

Outcome: Enhanced data processing capabilities, enabling more personalized customer experiences and data-driven decision-making.

5. NASA/JPL: Mars Rover

Challenge: NASA's Jet Propulsion Laboratory required a reliable platform to manage data from the Mars Rover.

Solution: Adopted AWS services like EC2 for compute, S3 for storage, and Glacier for long-term data archiving.

Outcome: Efficiently processed and stored vast amounts of scientific data, supporting ongoing Mars exploration missions.

6. Slack: Real-Time Messaging

Challenge: Slack needed a scalable and reliable infrastructure to support its real-time messaging platform.

Solution: Utilized AWS services including EC2 for scalable compute, RDS for managed databases, and Lambda for serverless operations.

Outcome: Achieved high availability and scalability, supporting millions of daily active users.

7. FINRA: Regulatory Compliance

Challenge: FINRA needed to process and store vast amounts of financial data to ensure regulatory compliance.

Solution: Used AWS services such as S3 for storage, Redshift for data warehousing, and EMR for big data processing.

Outcome: Improved data processing efficiency and regulatory compliance, handling trillions of records daily.

Chapter 8

Future Trends in Cloud Computing and AWS

1. Increased Adoption of Multi-Cloud Strategies:

Organizations are increasingly adopting multi-cloud strategies to leverage the strengths of different cloud providers, avoid vendor lock-in, and enhance resilience. AWS is likely to expand its interoperability with other cloud platforms and offer tools for seamless multi-cloud management.

2. Expansion of AI and Machine Learning Services:

The demand for AI and machine learning is growing rapidly. AWS will continue to enhance its suite of AI/ML services, like Amazon SageMaker, to support more advanced algorithms, easier model deployment, and broader integration with other AWS services.

3. Growth of Serverless Computing:

Serverless computing, exemplified by AWS Lambda, will become more prevalent due to its efficiency and cost-effectiveness. AWS is expected to expand serverless capabilities, support more programming languages, and improve integration with other AWS services.

4. Edge Computing and IoT:

With the rise of edge computing and the Internet of Things (IoT), AWS will likely introduce more services and enhancements for AWS IoT and AWS Greengrass. This will enable better processing and analysis of data closer to the data source, reducing latency and bandwidth usage.

5. Enhanced Security and Compliance:

As cybersecurity threats evolve, AWS will continue to innovate in security services. This includes improvements in identity and access management, encryption, threat detection, and compliance tools to meet stringent regulatory requirements.

6. Quantum Computing:

AWS is investing in quantum computing with Amazon Braket. Future trends will likely see significant advancements in making quantum computing more accessible and practical for solving complex problems.

7. Sustainability Initiatives:

Environmental sustainability is becoming a major focus. AWS will likely enhance its efforts in reducing carbon footprints, promoting green data centers, and providing tools for customers to manage their sustainability goals.

8. Enhanced Data Analytics and Big Data Solutions:

The need for advanced data analytics is growing. AWS will likely introduce more powerful big data tools and services for real-time analytics, data lakes, and data warehousing, making it easier to extract insights from large datasets.

9. Hybrid and Private Cloud Solutions:

As organizations seek flexibility, AWS will enhance hybrid cloud solutions like AWS Outposts and VMware Cloud on AWS, providing seamless integration between on-premises environments and the AWS cloud.

Chapter 9

Conclusion

In summary, internship training in cloud computing with AWS offers an invaluable, hands-on experience that equips interns with the practical skills necessary to thrive in the IT industry. Interns engage with real-world projects, leveraging a range of AWS services such as EC2, S3, Lambda, and RDS, which reinforces their theoretical knowledge through practical application.

Additionally, interns develop technical proficiency in cloud infrastructure management, learning to design scalable, resilient, and cost-effective solutions. Exposure to AWS best practices, including security, cost management, and operational excellence, ensures they can build efficient and secure cloud environments. The AWS Well-Architected Framework serves as a foundation for these best practices, guiding interns in creating robust solutions.

The completion of an AWS-focused internship significantly boosts employability, demonstrating to potential employers that the intern possesses a solid foundation in cloud computing and can contribute effectively to complex cloud projects. Moreover, interns gain awareness of future trends in cloud computing, such as serverless architecture, edge computing, and AI/ML integration, ensuring they remain prepared for upcoming technological advancements.

In conclusion, an internship in cloud computing with AWS is a transformative experience that bridges the gap between academic knowledge and industry practice. It provides interns with the skills, knowledge, and confidence needed to excel in the cloud computing domain, positioning them at the forefront of technological innovation as businesses increasingly migrate to the cloud.

Bibliography

- [1] AWS Automation Cookbook: Continuous Integration and Continuous Deployment using AWS services eBook : Swaraj, Nikit
- [2] <https://www.w3schools.com>
- [3] Geeks for Geeks, Geeks for Geeks, accessed July 13 , 2024, <https://geeksforgeeks.org>